

## What You'll Learn

- Learn what a Jupyter Notebook is and describe its role in modern data science, education, and reproducible research.
- Learn how front-end and kernel interact to run code and display outputs.
- Navigate the JupyterLab interface.
- Create and run code cells, including inline visualizations.
- Compose Markdown cells using headings, emphasis, lists, code blocks, images, and LaTeX equations to document analyses.
- How to work efficiently in JupyterLab.
- How you can troubleshoot common notebook issues.

# 1 Introduction to Jupyter Notebooks

A **Jupyter Notebook** (file extension `.ipynb`) is a single, shareable document that blends executable code, narrative text, equations, and rich media. A jupyter notebook runs in the browser and allows you to write, and execute, and annotate your thought process in styled Markdown. Because notebooks run on every major operating system and are recognised by research labs, startups, and Fortune-500 data teams alike, they have become the *lingua franca* of modern data science, machine learning, and technical education.

## 1 What is a Jupyter Notebook?

A **Jupyter Notebook** is a single, shareable document in your web browser where you can mix runnable code, plain-language notes, maths, images and charts in whatever order tells the clearest story — no special desktop software required. It has become the go-to format for teaching, data science and exploratory programming because anyone with Jupyter installed can open the file and replay every step exactly as you wrote it.

### 1.1 It's really just an ordinary text file

- **Human-readable “recipe card.”** A notebook is saved with the extension `.ipynb` and, under the hood, that file is plain text written in **JSON (JavaScript Object Notation)** — a simple way of organising data with curly braces and quotes, much like a digital spreadsheet made out of text.
- **Why JSON matters.** Because JSON is pure text, tools such as Git can show *exactly* what changed when you edit a notebook, line by line, just like they would for a normal document.
- **No coding knowledge required to understand that file exists.** You never have to look at the JSON directly, but it helps to know the file is nothing magical — you can email it, put it in Dropbox, or back it up like any other document.


### 1.2 Behind the scenes: two moving parts

Part	In plain English	What it does
<b>Front-end (the web page you see)</b>	The notebook editor that lives in your browser window.	Lets you create <i>cells</i> for code or text, click run, and view results instantly.
<b>Kernel (a tiny Python program running in the background)</b>	Think of it as your personal Python interpreter waiting for instructions.	Executes the code you type, then sends the answers— numbers, tables, plots—back to the front-end.

This “front-end plus kernel” design lets the same notebook speak to many languages (Python, R, Julia), simply by swapping kernels.

### 1.3 Jupyter Notebooks for Sharing Work

Because everything — inputs, outputs, notes and pictures—lives inside that one `.ipynb` file, anyone can double-click it on their own machine (or open it online with viewers such as **nbviewer**) and see precisely what you saw, right down to the plots. That makes notebooks perfect for homework, research papers, reports, or sharing a quick demo with a colleague.



## 2 A User-Friendly Tour of the Interface

When you first open JupyterLab, the screen is thoughtfully organized into several key areas:

### **Toolbar**

At the very top, you'll see buttons like **Run ▶**, **Stop ■**, **Restart ↺**, and menus labeled *File*, *Edit*, *View*, *Run*, *Kernel*, *Settings*, and *Help*. These are non-coders' shortcuts: instead of memorizing keyboard commands, you can click to run individual cells, restart your Python session, save your work, or export your notebook as a PDF. This visual interface ensures that you're never left guessing what button to press.

---

## Cell Gutter (Left Margin of Each Cell)

Next to each cell, you'll notice a narrow gray bar with labels like **In[ ]** or **In[3]**. This is known as the *cell gutter*. It serves two purposes:

1. Highlights which cell is currently selected.
2. Tracks the order in which you've executed cells. The number inside the brackets increments as you run new cells, helping you keep tabs on whether your code is updating in the expected sequence.

In other words, the gutter helps prevent confusion—no more wondering whether you forgot to run a certain cell.

---

## Sidebar & File Browser (Left Pane)

On the far left, the sidebar displays a vertical set of icons—by default, featuring a folder icon, among others. Click the folder icon to open the File Browser, which shows all files in your current project folder. Here you can:

- Double-click `.ipynb` notebooks, CSV, or script files to open them in new tabs.
- Upload new data files or create folders.
- Launch a terminal or see your running Python sessions.

This design lets you toggle between data, code, and visualizations without leaving the notebook—no app switching required.

---

## Main Work Area (Center Workspace)

This is where the real magic happens: your notebook lives here. You'll create **cells**—blocks of code or formatted text—and run them one by one. Each cell can include Python code that executes immediately, or markdown text that renders into readable documentation. The interface feels like a powerful word processor designed especially for coding and explanation.

---

## Status Bar (Bottom)

Along the bottom lies a thin bar showing important contextual info:

- **Kernel name** ("Python3.13") – tells you which programming language/environment is running.
- **CPU/Memory usage** – helpful when working with large data or complex models.
- **Connection indicators** – shows if your Python kernel is busy or disconnected.

This is your constant reassurance that code is running where you expect, and that your session is actively executing commands.

## Why These Tools Matter

- **Clickable toolbar**: No one needs to remember `Shift + Enter`—just click Run and focus on ideas, not key combinations.
  - **Cell gutter numbering**: Helps you identify which cells haven't been run yet, making debugging easier.
  - **Sidebar navigation**: Lets you manage data files, switch between notebooks or open terminals—all in a single window.
  - **Status bar feedback**: Gives you immediate awareness of the system you're working on and the notebook's performance.
- 



## Putting It All Together

1. **Click the file in the sidebar** to open the notebook.
  2. Select a **cell** and inspect it—notice gutter highlighting.
  3. Click **Run ▶** to execute it; output appears right below the cell.
  4. Monitor the **status bar** to confirm the kernel is active and working.
  5. Use the **toolbar** and **menus** to restart the kernel, save, or export work as needed.
- 

## Why this setup is the gold standard in 2025

- **Transparent feedback** lets you see code execution and results side by side.
  - **Integrated data navigation** enables workflows that mix coding, exploration, and documentation—all without switching apps.
  - **Reproducibility** is built in: the order of cell execution, environment, and outputs are retained in a single shareable file.
  - **Industry-ready workflow**: as tools like VS Code's Jupyter integration, Cursor, GitHub Codespaces, and WinSurf adopt the same interface patterns, every skill you build here transfers directly into modern professional environments.
- 

## Tips for Beginners

## Tips for Beginners

Tip	Why It Helps
<b>Hover over icons</b> in the sidebar to see tooltips (like “File Browser”).	Useful when you're just learning what each panel does.
<b>Toggle the sidebar</b> with <code>Ctrl + B</code> (Cmd on Mac).	Allows for a distraction-free view while reading or presenting.
<b>Restart kernel and rerun all</b> from <i>Run</i> menu if outputs get messy.	Ensures your notebook behaves predictably from a clean start.

## 3 Cells: the Building Blocks

### 3.1 Code Cells

Paste the snippet below into the first cell and press **Shift+Enter** (or click *Run*):

```
# My first calculation
3 + 4
```

The number 7\* appears right below.\* Congratulations, you just executed code inside the browser! Keyboard shortcuts like **Shift+Enter** (run & advance) and **Ctrl+Enter** (run & stay) are the fastest way to iterate.

Add a second code cell and test inline plotting:

```
%matplotlib inline
import matplotlib.pyplot as plt
plt.plot([1, 2, 3], [4, 6, 5])
plt.title("Inline plot demo")
```

With the magic `%matplotlib inline`, the figure renders *inside* the notebook instead of a pop-up window, keeping narrative, code, and visuals together.

### 3.2 Markdown Cells

Switch a fresh cell to **Markdown** (press **M** in command mode or choose *Cell → Cell Type → Markdown*). Type:

```
## Exploratory Analysis
```

We expect **price** to rise as *size* increases.

- Data source: ``data/housing.csv``
- Sample: 500 rows

The modelled relationship is  $y = \beta_0 + \beta_1 x + \epsilon$ .

Run the cell. Text is rendered with headings, bold and italics, an inline code font, bullet list, and a LaTeX equation—all without leaving the notebook.

## 4 Why Markdown Matters in Data Science

- **Documentation & storytelling** – analyses become *explainable* when every decision is written next to the code that produced it. Data scientists routinely submit notebooks as final reports.
- **Reproducibility** – lab notes, dataset descriptors, and parameter choices live beside the results, ensuring peers can replicate your work months later.
- **Collaboration** – version-control diffs are human-readable; a teammate can review narrative and code in one pull request.

The 2024 Stack Overflow survey lists Markdown as the most admired documentation tool among 84 % of developers, showing its cross-industry relevance.

### 4.1 Essential Markdown Cheatsheet

Syntax	Renders as
# Heading 1 → ##### Heading 6	Six heading sizes
<b>**bold**</b> , <i>*italic*</i> , <code>`code`</code>	Emphasis & inline code
- / * / 1.	Bulleted or numbered lists
python ...	Fenced code block with optional language
> Blockquote	Quoted text
---	Horizontal rule
![alt text](image.png)	Embedded image
$E = mc^2$	Inline LaTeX maths
$\int_{-\infty}^{\infty} e^{-x^2} \, dx$	Display-mode equation

The following screenshot illustrates both the raw Markdown and the rendered output side-by-side.



## 5 Beyond the Basics

### 5.1 Interactive Widgets

Libraries such as **ipywidgets** add sliders, dropdowns, and play buttons so you can poke at parameters live. A one-liner:

```
from ipywidgets import interact
@interact(k=(-2, 2, 0.1))
def line(k=1):
    plt.figure()
    plt.plot([0,1], [0,k])
```

Move the slider and the plot updates instantly—great for teaching and exploratory analysis.

### 5.2 Keyboard Shortcuts You’ll Use Daily

Action	Command Mode	Edit Mode
Run cell & advance	<b>Shift+Enter</b>	same
Run cell, keep focus	<b>Ctrl+Enter</b>	same
Add <b>Code</b> cell below	<b>B</b>	<b>Esc→B</b>
Add <b>Markdown</b> cell	<b>Esc→M</b>	—
Delete cell	<b>Esc→DD</b>	—

Master these shortcuts to cut your mouse mileage in half.

## 6 How Notebooks Fit into the Bigger Toolchain

Many modern editors embed notebook tech so you can pivot between script-centric and cell-centric workflows without exporting files. Learning the plain-vanilla notebook today transfers directly to these tools when we move into these tools later. VS Code, for example, shows notebooks with the same run buttons and Markdown rendering, plus side-by-side Git diffs and Copilot suggestions.

## 7 Troubleshooting Checklist

Issue	Quick check
Output appears in the wrong place	Ensure cells run top-to-bottom; restart kernel ☛ if execution count jumps.
<code>ModuleNotFoundError</code> despite having installed a package	Verify you installed it <b>inside</b> the active virtual environment & kernel.
Notebook uses huge RAM	Clear outputs (Menu → Kernel → Restart & Clear Output) before sharing or committing to a code repository (more on this later)